

## ✓ Project-2 : Spam Detection Using TensorFlow(Python)

### ✓ 0. Import the project dependencies

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import tensorflow as tf
import nltk
import string
from nltk.corpus import stopwords
import sklearn
from wordcloud import WordCloud
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.callbacks import EarlyStopping

import warnings
warnings.filterwarnings('ignore')
```

Start coding or [generate](#) with AI.

### 1. Problem Statement

The Objective is to build a Machine Learning model to detect spam emails.

- Prepare the dataset for machine learning including data cleaning, removing missing values and splitting the data into training and test sets.
- Building a TensorFlow Model to detect spam emails.
- Evaluate the performance of the model on the test dataset
- Analyze the model coefficients to understand the importance of the different features.
- Build a report that describes the steps you took to complete the task. the result of our analysis and conclusions
- The code that i used to build also describe it.

### ✓ 2. Relevent Dataset

```
data = pd.read_csv("/content/spam_ham_dataset.csv")
data.head()
```

	Unnamed: 0	label	text	label_num	
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...	0	
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n( see...	0	
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...	0	
3	4685	spam	Subject: photoshop , windows , office . cheap ...	1	
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...	0	

Next steps:

[Generate code with data](#)[View recommended plots](#)

### label vs Unnamed: 0

```
# @title label vs Unnamed: 0
```


```
from matplotlib import pyplot as plt
import seaborn as sns
figsize = (12, 1.2 * len(data['label'].unique()))
plt.figure(figsize=figsize)
sns.violinplot(data, x='Unnamed: 0', y='label', inner='box', palette='Dark2')
sns.despine(top=True, right=True, bottom=True, left=True)
```



McAfee WebAdvisor




Your download's being scanned.  
We'll let you know if there's an issue.



### 3. Data Exploration

#### Visualizing the data distribution

```
sns.countplot(x='label', data = data)
plt.show()
```



#### Dataset Balancing

```
email_msg = data[data.label_num == 0]
spam_msg = data[data.label_num == 1]
email_msg = email_msg.sample(n=len(spam_msg), random_state=42)
```

#### Let's Visualize the balanced dataset


```
balanced_data = pd.concat([email_msg, spam_msg], ignore_index=True)
plt.figure(figsize=(8 , 6))
sns.countplot(data = balanced_data, x = 'label')
plt.title('Distribution of Email and Spam messages after downsampling')
plt.xlabel('Messages Class')
plt.show()
```



McAfee WebAdvisor



Your download's being scanned.  
We'll let you know if there's an issue.



#### 4. Data Preprocessing

##### Let's Preprocess the text dataset

```

nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True

balanced_data['text'] = balanced_data['text'].str.replace('Subject', '')
punctuations_list = string.punctuation
def rem_punch(text):
    temp = str.maketrans('', '', punctuations_list)
    return text.translate(temp)

balanced_data['text'] = balanced_data['text'].apply(lambda x: rem_punch(x))

def rem_stopwords(text):
    stop_words = stopwords.words('english')

    imp_words = []
    for word in str(text).split():
        word = word.lower()

        if word not in stop_words:
            imp_words.append(word)
    output = " ".join(imp_words)

    return output

balanced_data['text'] = balanced_data['text'].apply(lambda text: rem_stopwords(text))
balanced_data.head()

```

Unnamed: 0	label	text	label_num
0	3444 ham	conoco big cowboy darren sure help know else a...	0
1	2982 ham	feb 01 prod sale teco gas processing sale deal...	0
2	2711 ham	california energy crisis california power cr...	0
3	3116 ham	nom actual volume april 23 rd agree eileen pon...	0
4	1314 ham	easttrans nomination changes effective 8 2 00 p...	0

Next steps: [Generate code with balanced\\_data](#)

[View recommended plots](#)



McAfee WebAdvisor

Your download's being scanned.  
We'll let you know if there's an issue.

- Let's visualize the text dataset using wordcloud

```
def plot_words(data, typ):
    sms_corpus = " ".join(data['text'])


    plt.figure(figsize=(7, 7))

    wc = WordCloud(
        background_color='black',
        max_words = 100,
        width = 800,
        height = 400,
        collocations = False
    ).generate(sms_corpus)


    plt.imshow(wc, interpolation='bilinear')
    plt.title(f'WordCloud for {typ} email', fontsize=15)

plot_words(balanced_data[balanced_data['label_num'] == 0], typ='not-spam')
plot_words(balanced_data[balanced_data['label_num'] == 1], typ='spam')
```

### WordCloud for not-spam email



### WordCloud for spam email



- Let's split the dataset into train and test sets

```
X_train, X_test, Y_train, Y_test = train_test_split(
    balanced_data['text'],
    balanced_data['label_num'],
    test_size=0.2,
    random_state=42
)
```

- Tokenization and Padding



 **McAfee** WebAdvisor

Your download's being scanned.  
We'll let you know if there's an issue.



```

tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train)
train_sequences = tokenizer.texts_to_sequences(X_train)
test_sequences = tokenizer.texts_to_sequences(X_test)
max_len = 100

train_sequences = pad_sequences(train_sequences,
                                maxlen = max_len,
                                padding = 'post',
                                truncating = 'post')

test_sequences = pad_sequences(test_sequences,
                                maxlen = max_len,
                                padding='post',
                                truncating = 'post')

```

## 5. Machine Learning Model

```

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Embedding(input_dim = len(tokenizer.word_index) + 1, output_dim=32, input_length=100))
model.add(tf.keras.layers.LSTM(16))
model.add(tf.keras.layers.Dense(32, activation='relu'))
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 100, 32)	1274912
lstm (LSTM)	(None, 16)	3136
dense (Dense)	(None, 32)	544
dense_1 (Dense)	(None, 1)	33

Total params: 1278625 (4.88 MB)  
Trainable params: 1278625 (4.88 MB)  
Non-trainable params: 0 (0.00 Byte)

```

model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics= ['accuracy']
)

history = model.fit(train_sequences, Y_train, epochs=20, batch_size=32, validation_data=(test_sequences, Y_test))

```

Epoch 1/20

75/75 [=====] - 16s 169ms/step - loss: 0.6910 - accuracy: 0.5334 - val\_loss: 0.6818 - val\_accuracy: 0.5883

Epoch 2/20

75/75 [=====] - 10s 133ms/step - loss: 0.3089 - accuracy: 0.9012 - val\_loss: 0.1548 - val\_accuracy: 0.9617

Epoch 3/20

75/75 [=====] - 8s 114ms/step - loss: 0.1284 - accuracy: 0.9708 - val\_loss: 0.1615 - val\_accuracy: 0.9617

Epoch 4/20

75/75 [=====] - 7s 96ms/step - loss: 0.1063 - accuracy: 0.9771 - val\_loss: 0.1412 - val\_accuracy: 0.9683

Epoch 5/20

75/75 [=====] - 7s 93ms/step - loss: 0.0984 - accuracy: 0.9791 - val\_loss: 0.1929 - val\_accuracy: 0.9567

Epoch 6/20

75/75 [=====] - 5s 66ms/step - loss: 0.1029 - accuracy: 0.9783 - val\_loss: 0.1803 - val\_accuracy: 0.9467

Epoch 7/20

75/75 [=====] - 5s 66ms/step - loss: 0.1177 - accuracy: 0.9700 - val\_loss: 0.0837 - val\_accuracy: 0.9800

Epoch 8/20

75/75 [=====] - 4s 51ms/step - loss: 0.0576 - accuracy: 0.9892 - val\_loss: 0.0924 - val\_accuracy: 0.9783

Epoch 9/20

75/75 [=====] - 3s 42ms/step - loss: 0.0339 - accuracy: 0.9942 - val\_loss: 0.1073 - val\_accuracy: 0.9767

Epoch 10/20

75/75 [=====] - 4s 49ms/step - loss: 0.0284 - accuracy: 0.9950 - val\_loss: 0.1065 - val\_accuracy: 0.9783

Epoch 11/20

75/75 [=====] - 3s 39ms/step - loss: 0.0262 - accuracy: 0.9958 - val\_loss: 0.1290 - val\_accuracy: 0.9750

Epoch 12/20

75/75 [=====] - 3s 41ms/step - loss: 0.0232 - accuracy: 0.9962 - val\_loss: 0.1383 - val\_accuracy: 0.9733

Epoch 13/20

75/75 [=====] - 3s 39ms/step - loss: 0.0202 - accuracy: 0.9967 - val\_loss: 0.1249 - val\_accuracy: 0.9767

Epoch 14/20

75/75 [=====] - 4s 51ms/step - loss: 0.0229 - accuracy: 0.9967 - val\_loss: 0.1249 - val\_accuracy: 0.9767


Epoch 15/20

75/75 [=====] - 4s 53ms/step - loss: 0.0210 - accuracy: 0.9967 - val\_loss: 0.1249 - val\_accuracy: 0.9767

Epoch 16/20


75/75 [=====] - 2s 26ms/step - loss: 0.0193 - accuracy: 0.9967 - val\_loss: 0.1249 - val\_accuracy: 0.9767

Epoch 17/20



McAfee

WebAdvisor



Your download's being scanned.

We'll let you know if there's an issue.

```

75/75 [=====] - 2s 25ms/step - loss: 0.0201 - accuracy: 0.9971 - val_loss: 0.1357 - val_accuracy: 0.9767
Epoch 18/20
75/75 [=====] - 2s 28ms/step - loss: 0.0174 - accuracy: 0.9975 - val_loss: 0.1316 - val_accuracy: 0.9767
Epoch 19/20
75/75 [=====] - 2s 21ms/step - loss: 0.0175 - accuracy: 0.9975 - val_loss: 0.1296 - val_accuracy: 0.9767
Epoch 20/20
75/75 [=====] - 2s 27ms/step - loss: 0.0173 - accuracy: 0.9975 - val_loss: 0.1301 - val_accuracy: 0.9767

```

## 6. Model Evaluation

```

test_loss, test_accuracy = model.evaluate(test_sequences, Y_test)
print("Test Loss :", test_loss)
print("Test Accuracy :", test_accuracy)

```

```

19/19 [=====] - 0s 4ms/step - loss: 0.1301 - accuracy: 0.9767
Test Loss : 0.13006488978862762
Test Accuracy : 0.9766666889190674


```

## Visualizing the model's test accuracy

```

plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()

```



McAfee WebAdvisor



Your download's being scanned.  
We'll let you know if there's an issue.